

Lab 5: Regression & Neural Networks

519021910702 游克垚

1 Exercise 1: Linear Regression

1.1 Implementation

My implementation is as follows:

```
1 import numpy as np
2
3 if __name__ == "__main__":
4     X = np.load('Data1_X.npy')
5     Y = np.load('Data1_Y.npy')
6
7     X = np.matrix(X)
8     Y = np.matrix(Y)
9
10    Y_hat = X * (X.T * X).I * X.T * Y
```

1.2 Result

The result of \hat{Y}^T is as follows:

```
1 [[ 90.19765634  47.67686144  91.65938765 120.74673813 115.32784698
2    99.5235342  101.6811349  73.64256436 118.02996154 101.22977512
3    53.90993768  80.72800342 100.06768951  84.87342942 108.10549645
4    84.28833524  87.50086209  87.23821599  82.31093561 104.35109097
5    69.9964474  97.5290999 113.73176903  83.13355445 125.47193094
6    91.16712382  87.03894768  94.06534565  72.53264476 105.78014471
7   109.47621827  93.9088358 110.22388501  83.04389589 105.0040709
8    64.34001661  58.32922243  90.39147089 113.32483653  62.51171231
9    86.12664693  78.76219229  58.93439827  42.12567826 126.62571301
10   106.33033386  46.80686218  64.43015537  60.12018271  58.64341279
11   103.05922202 108.34444014  59.46080212 136.30383353  86.79373439
12    81.59990756  91.86133364 100.73337199  86.9771913  70.95011886
13    75.72726683  97.88306054  34.9951482 100.82833137 109.88897963
14    92.81345368  89.46905032 136.37900461  65.90720482  82.27029322
15   111.67846566  46.88349888  69.84761845  95.91483137  62.0903558
16    74.21588583  79.1009826  71.19647357 127.98576243  59.55125149
```

```

17 94.22580938 82.5433742 74.60546716 87.18355727 99.21624982
18 91.55587898 114.7186236 107.22598751 89.13747846 112.2219674
19 91.2314803 104.22842129 102.5151316 109.87026128 117.48132199
20 78.98080158 50.16010039 85.87762567 85.88482494 106.62073286]]

```

Under the geometric interpretation of linear regression, our goal is to find a point \hat{Y} in the X plane which is the nearest to the point Y . So vector $Y - \hat{Y}$ must be perpendicular to the X plane. So

$$\begin{aligned}
 X^T(Y - \hat{Y}) &= 0 \\
 X^T(Y - X\hat{\theta}) &= 0 \\
 \hat{\theta} &= (X^T X)^{-1} X^T Y
 \end{aligned}$$

So $\hat{Y} = X\hat{\theta} = X(X^T X)^{-1} X^T Y$.

2 Exercise 2: Logistic Regression

2.1 Implementation

My implementation is as follows:

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from utils import plot_decision_boundary
4
5 def CELoss_binary(X, y, theta):
6     z = X @ theta
7     y_hat = 1 / (1 + np.exp(-z))
8     return - (y.T @ np.log(y_hat + 0.000001) + (1-y).T @ np.log(1-y_hat + 0.000001)) /
           X.shape[0]
9
10 def precision(y_true, y_predict):
11     return np.sum(y_true == y_predict) / y_true.shape[0]
12
13 def predict(X, theta, threshold):
14     z = X @ theta
15     y_hat = 1 / (1+np.exp(-z))
16     y_predict = np.zeros(y_hat.shape)
17     y_predict[y_hat > threshold] = 1
18     return y_predict
19
20 def gradient(X, y, theta):
21     z = X @ theta
22     y_hat = 1 / (1 + np.exp(-z))
23     return X.T @ (y_hat-y) / X.shape[0]
24
25 if __name__ == "__main__":

```

```

26     X_all, y_all = np.load('Data2_X.npy'), np.load('Data2_Y.npy')
27
28     b = np.ones((X_all.shape[0], 1))
29     X_all = np.c_[X_all, b]
30
31     for k in range(3):
32         theta = np.ones((X_all.shape[1],))
33
34         lr = 0.01
35         iteration_num = 1000
36         threshold = 0.2 + k * 0.3
37
38         loss = []
39         prec = []
40
41         for i in range(iteration_num):
42             # change algorithm
43             index = np.random.choice(np.arange(X_all.shape[0]), size=100, replace=False
44                                     )
45             X = X_all[index]
46             y = y_all[index]
47             theta -= lr * gradient(X, y, theta)
48
49             y_predict = predict(X_all, theta, threshold)
50
51             loss.append(CELoss_binary(X_all, y_all, theta))
52             prec.append(precision(y_all, y_predict))
53
54             # plt.plot([i for i in range(iteration_num)], loss, label=str(threshold))
55             # plt.plot([i for i in range(iteration_num)], prec, label=str(threshold))
56
57             plot_decision_boundary(X_all, y_all, lambda x : predict(np.c_[x, np.ones((x.
58                                     shape[0], 1))], theta, threshold))
59             plt.show()
60
61         # plt.legend()
62
63         # plt.xlabel('iteration times')
64         # plt.ylabel('precision')
65         # plt.ylabel('binary cross entropy loss')
66
67         # plt.show()

```

When counting `CELoss_binary`, I add a small number to \hat{y} to avoid $np.log(0)$.

2.2 Binary cross entropy loss

The graphs of binary cross entropy loss against the number of iterations using stochastic gradient descent, mini-batch gradient descent and (batch) gradient descent respectively under 3 different learning rates are as follows:

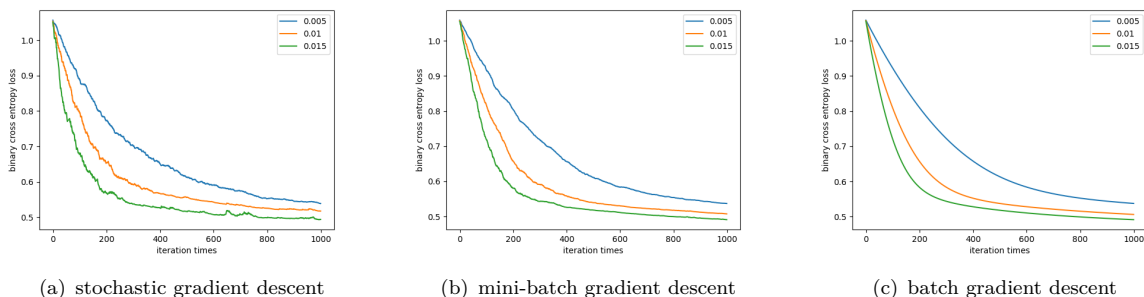


Fig 2-1 binary cross entropy loss under 3 different learning rates

If the learning rate is larger, the binary cross entropy loss descends quicker. And the binary cross entropy loss fluctuates most when using stochastic gradient descent, less when using mini-batch gradient descent and least when using batch gradient descent. The stochastic gradient descent is also the fastest, followed by mini-batch gradient descent, and the batch gradient descent is the slowest.

The graphs of binary cross entropy loss against the number of iterations using stochastic gradient descent, mini-batch gradient descent and (batch) gradient descent respectively under 3 different values of the threshold are as follows:

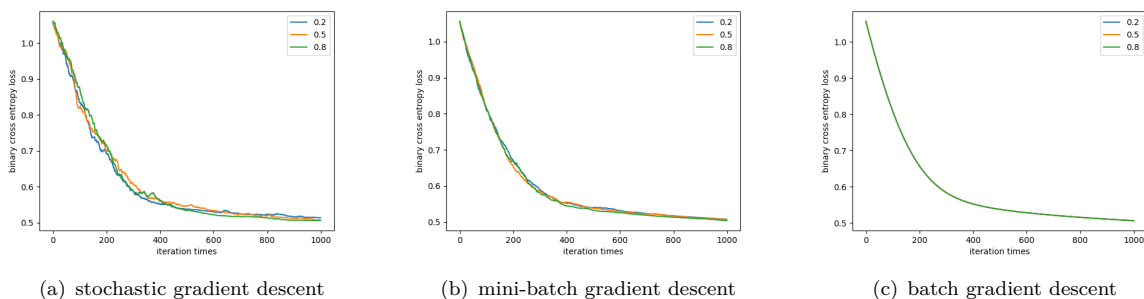


Fig 2-2 binary cross entropy loss under 3 different values of the threshold

The value of threshold doesn't effect binary cross entropy loss much because it only works when predicting. The effects of types of gradients are the same.

2.3 Precision

The graphs of precision against the number of iterations using stochastic gradient descent, mini-batch gradient descent and (batch) gradient descent respectively under 3 different learning rates are as follows:

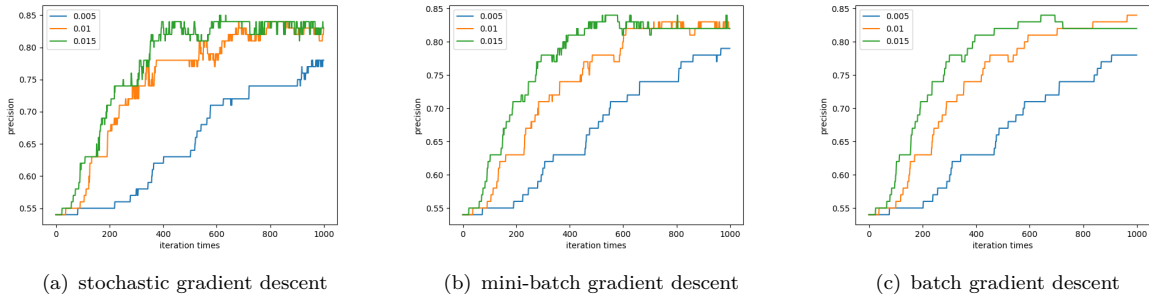


Fig 2-3 precision under 3 different learning rates

Similar to the graphs above, the learning rate influences the speed of learning. But when the learning rate is too large, it is difficult to converge and the precision is low. The effects of types of gradients are the same.

The graphs of precision under 3 different values of the threshold are as follows:

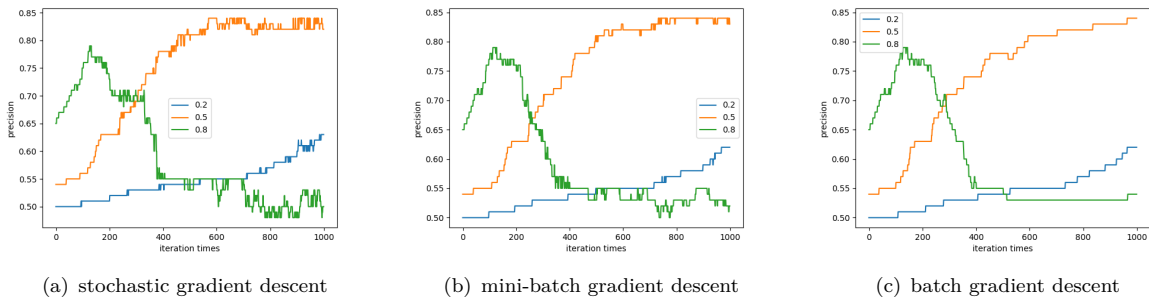


Fig 2-4 precision under 3 different values of the threshold

The value of the threshold is important to precision because it decides the decision boundary. The effects of types of gradients are the same.

2.4 Decision boundary

The graphs of decision boundary of predictions with different values of threshold are as follows:

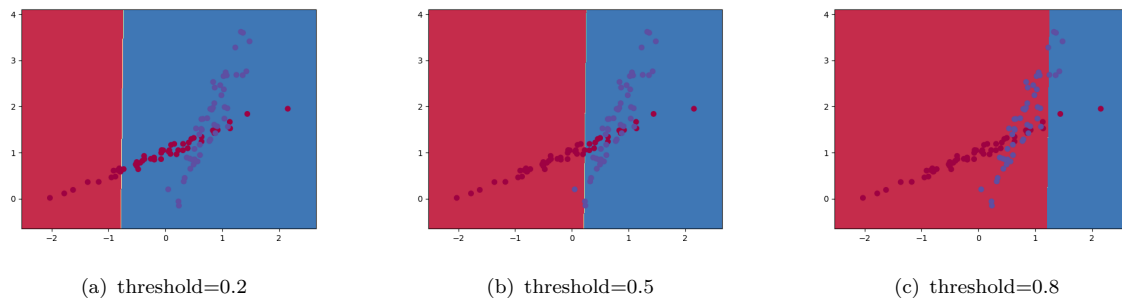


Fig 2-5 decision boundary of predictions with different values of threshold

We can find that if the value of threshold is smaller, the decision boundary moves more in the

negative direction of the x-axis. But with the number of iteration increasing, they will gradually approach to each other.

3 Exercise 3: L1/L2 Regularization

3.1 Ridge regression

My implementation is as follows:

```

1 def ridge_regression(x, y, lamda):
2     # Normalize data.
3     x = (x - x.mean(axis=-1, keepdims=True)) / x.std(axis=-1, keepdims=True)
4     from sklearn.linear_model import Ridge
5     ridge = Ridge(alpha=lamda)
6     ridge.fit(x, y)
7     y_pred = ridge.predict(x) # predicted labels of size (n_samples, )
8     intercept = ridge.intercept_ # b of size ( )
9     coef = ridge.coef_ # theta of size (n_dims, )
10    return y_pred, intercept, coef

```

The result of ridge regression is as follows:

```

1 0.03876948356628418 1.042289124220179 0.0
2 0.0012583732604980469 1.1052304919020004 0.0
3 0.0005276203155517578 1.5737740179631936 0.0
4 0.0006253719329833984 1.6765680529544285 0.0
5 0.000644683837890625 1.732565285766469 0.0
6 0.0005466938018798828 7.331330534314156 0.0

```

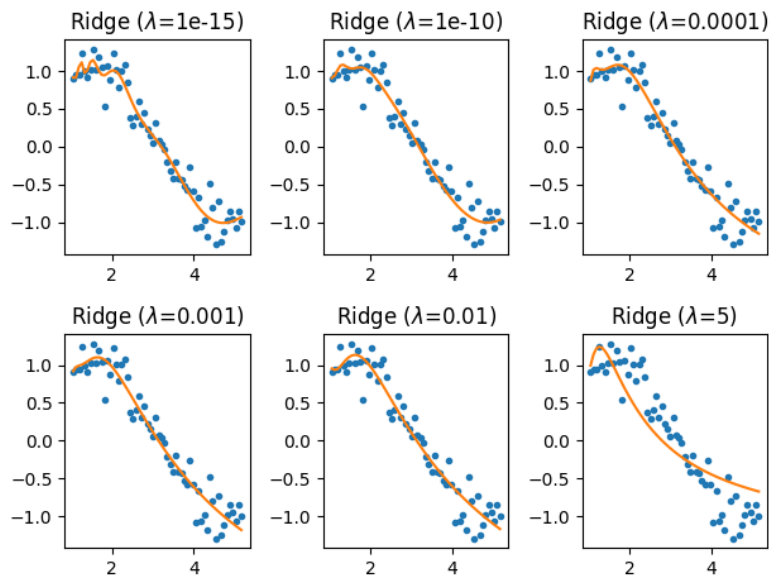


Fig 3-6 result of ridge regression

3.2 Lasso regression

My implementation is as follows:

```

1 def lasso_regression(x, y, lamda):
2     # Normalize data.
3     x = (x - x.mean(axis=-1, keepdims=True)) / x.std(axis=-1, keepdims=True)
4
5     from sklearn.linear_model import Lasso
6
7     lasso = Lasso(alpha=lamda, max_iter=1000000)
8     lasso.fit(x, y)
9
10    y_pred = lasso.predict(x) # predicted labels of size (n_samples, )
11    intercept = lasso.intercept_ # b of size ( )
12    coef = lasso.coef_ # theta of size (n_dims, )
13
14    return y_pred, intercept, coef

```

The result of lasso regression is as follows:

```

1 1.448115348815918 1.333009473823675 0.0
2 0.4232501983642578 1.5347481924504889 68.75
3 0.003134012222290039 1.7176378773851337 75.0
4 0.0009174346923828125 1.8505307867417076 81.25
5 0.0005638599395751953 3.134818651748969 87.5
6 0.000553131103515625 40.4165451776714 100.0

```

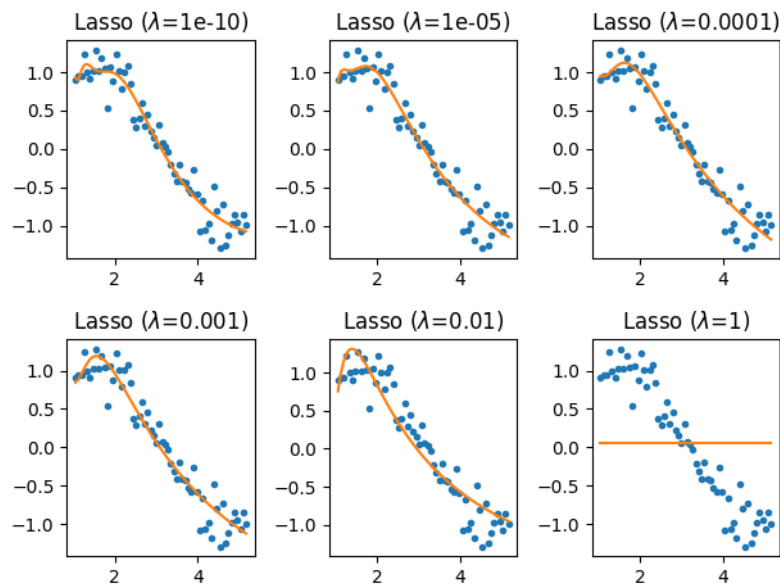


Fig 3-7 result of lasso regression

3.3 Discussion

We can find ridge regression spends less time while lasso regression spends more time. So ridge regression is less computationally expensive. And lasso regression tends to create a sparser output. Generalization refers a model adjusting itself based on the training data. Regularization refers a model reducing coefficients in learning and preventing overfitting. So regularization is a method to increase generalization.

4 Exercise 4: Two-layer Perceptron Network

4.1 Implementation

My implementation is as follows:

```

1  class Relu(Layer):
2      def __init__(self, name):
3          super(Relu, self).__init__(name)
4
5      def forward(self, input):
6          self._saved_for_backward(input)
7          input[input <= 0] = 0
8
9          return input
10
11     def backward(self, grad_output):
12         grad_output[self._saved_tensor <= 0] = 0
13
14         return grad_output
15
16 class Linear(Layer):
17     def __init__(self, name, in_num, out_num, init_std):
18         super(Linear, self).__init__(name, trainable=True)
19
20         self.in_num = in_num
21         self.out_num = out_num
22         self.W = np.random.randn(in_num, out_num) * init_std
23         self.b = np.zeros(out_num)
24
25         self.grad_W = np.zeros((in_num, out_num))
26         self.grad_b = np.zeros(out_num)
27
28         self.diff_W = np.zeros((in_num, out_num))
29         self.diff_b = np.zeros(out_num)
30
31     def forward(self, input):
32         self._saved_for_backward(input)
33

```



```
34         return input.dot(self.W) + self.b
35
36     def backward(self, grad_output):
37         self.grad_W = self._saved_tensor.T @ grad_output
38         self.grad_b = grad_output.sum(axis=0)
39
40         return grad_output @ self.W.T
41
42     def update(self, config):
43         mm = config['momentum']
44         lr = config['learning_rate']
45
46         self.W -= lr * self.grad_W
47         self.b -= lr * self.grad_b
48
49 class Network(object):
50     def __init__(self):
51         self.layer_list = []
52         self.params = []
53         self.num_layers = 0
54
55     def add(self, layer):
56         self.layer_list.append(layer)
57         self.num_layers += 1
58
59     def forward(self, input):
60         output = input
61         for i in range(self.num_layers):
62             output = self.layer_list[i].forward(output)
63
64         return output
65
66     def backward(self, grad_output):
67         grad = grad_output
68         for i in range(self.num_layers-1, -1, -1):
69             grad = self.layer_list[i].backward(grad)
70
71         return grad
72
73     def update(self, config):
74         for i in range(self.num_layers):
75             if self.layer_list[i].trainable:
76                 self.layer_list[i].update(config)
77
78     def predict(self, input):
79         y_pred = self.forward(input).argmax(axis=-1)
80
81         return y_pred
```

```

82
83 class EuclideanLoss:
84     def __init__(self, name):
85         self.name = name
86
87     def forward(self, input, target):
88         return ((target - input) ** 2).mean(axis=0).sum() / 2.
89
90     def backward(self, input, target):
91         return (input - target)

```

4.2 Result

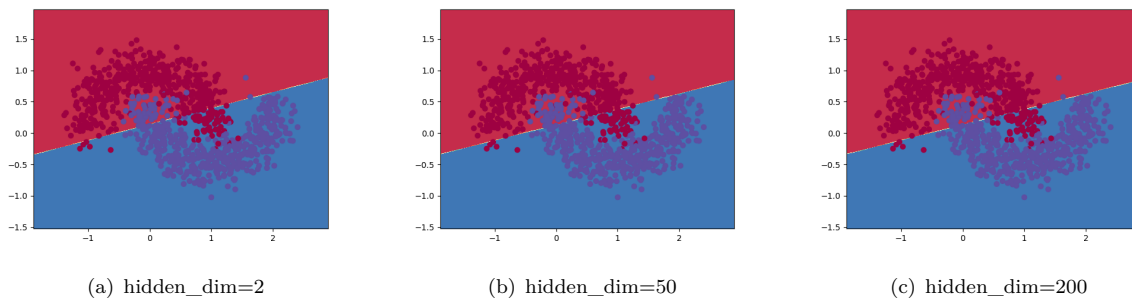


Fig 4-8 result of two-layer perceptron network

The result of two-layer perceptron network having different numbers of hidden neurons is the same. But if the number of hidden neurons is smaller, it is more likely to produce a failure as below.

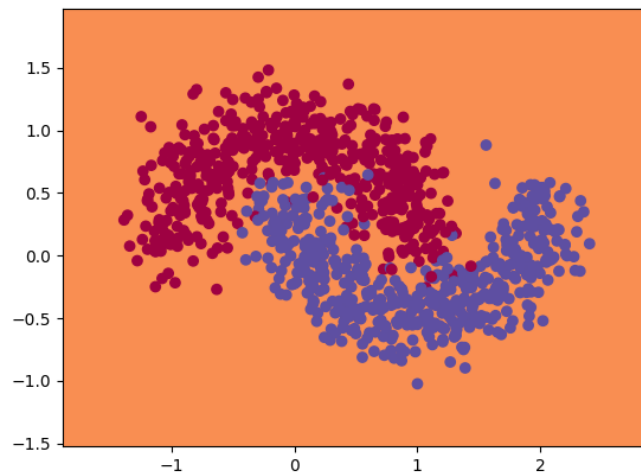


Fig 4-9 a failure of training